

Policy Iteration

Policy Game	Iteration Score
0	770
1	810
2	410
3	780
4	580
5	290
6	480
7	420
8	560
9	460
10	320
11	460
12	500
13	480
14	570
15	720
16	490
17	400
18	400
19	730
Average: 531.5	

I tested the PolicyIterationAgent against NullGhosts in 20 games. The scores ranged from (290 to 810), with an average of 531.5 points. The scores are different in each game because the test runner adds a bit of randomness to Pac-Man's moves, while the learned policy itself comes from a deterministic policy iteration process with stationary ghosts.